

multiple-regression-analysis.R

Sharad Deshmukh

2022-03-01

```
D=read.table("C:\\Users\\Sharad Deshmukh\\Desktop\\MSC=SEM-
II\\practical\\NationalFootballLeague.txt",header=TRUE,sep="\t",dec=".")
##Data preproccing
summary(D) ##No missing value

##           X           y           x1           x2           x3
## Min.      : 1.00    Min.      : 0.000    Min.      :1416    Min.      :1414    Min.
:35.10
## 1st Qu.: 7.75    1st Qu.: 4.000    1st Qu.:1896    1st Qu.:1714    1st
Qu.:37.38
## Median :14.50    Median : 6.500    Median :2111    Median :2106    Median
:38.85
## Mean     :14.50    Mean      : 6.964    Mean      :2110    Mean      :2127    Mean
:38.64
## 3rd Qu.:21.25    3rd Qu.:10.000    3rd Qu.:2303    3rd Qu.:2474    3rd
Qu.:39.70
## Max.      :28.00    Max.      :13.000    Max.      :2971    Max.      :2929    Max.
:42.30
##           x4           x5           x6           x7
## Min.      :38.10    Min.      :-22.00    Min.      : 576.0    Min.      :43.80
## 1st Qu.:52.42    1st Qu.: -5.75    1st Qu.: 710.5    1st Qu.:54.77
## Median :57.70    Median :  1.00    Median : 787.5    Median :58.65
## Mean     :59.40    Mean      :  0.00    Mean      : 789.9    Mean     :58.16
## 3rd Qu.:68.80    3rd Qu.:  6.25    3rd Qu.: 869.8    3rd Qu.:61.10
## Max.      :78.30    Max.      :19.00    Max.      :1037.0    Max.      :67.50
##           x8           x9
## Min.      :1457    Min.      :1575
## 1st Qu.:1848    1st Qu.:1913
## Median :2050    Median :2101
## Mean     :2110    Mean      :2128
## 3rd Qu.:2320    3rd Qu.:2328
## Max.      :2876    Max.      :2670

##reducing the dimension
dim(D)

## [1] 28 11

names(D)

## [1] "X"  "y"  "x1" "x2" "x3" "x4" "x5" "x6" "x7" "x8" "x9"

str(D)
```

```
## 'data.frame': 28 obs. of 11 variables:
## $ X : int 1 2 3 4 5 6 7 8 9 10 ...
## $ y : int 10 11 11 13 10 11 10 11 4 2 ...
## $ x1: int 2113 2003 2957 2285 2971 2309 2528 2147 1689 2566 ...
## $ x2: int 1985 2855 1737 2905 1666 2927 2341 2737 1414 1838 ...
## $ x3: num 38.9 38.8 40.1 41.6 39.2 39.7 38.1 37 42.1 42.3 ...
## $ x4: num 64.7 61.3 60 45.3 53.8 74.1 65.4 78.3 47.6 54.2 ...
## $ x5: int 4 3 14 -4 15 8 12 -1 -3 -1 ...
## $ x6: int 868 615 914 957 836 786 754 761 714 797 ...
## $ x7: num 59.7 55 65.6 61.4 66.1 61 66.1 58 57 58.9 ...
## $ x8: int 2205 2096 1847 1903 1457 1848 1564 1821 2577 2476 ...
## $ x9: int 1917 1575 2175 2476 1866 2339 2092 1909 2001 2254 ...

D=D[,-1]
wineClasses <- factor(D[,1])

#plot(main="Three Different
Cultivars",D[,2],D[,3],D[,4],D[,5],D[,6],D[,7],D[,8],D[,9],D[,10], col =
wineClasses)

##PCA
dimPCA <- prcomp(scale(D[, -1]))
#Step 3: Choose the principal components with highest variances
#Now the 13 features has reduced to only 2 new Principal Components These are
not 2 of those 13, but 2 new components
summary(dimPCA)

## Importance of components:
##
## Standard deviation      PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Proportion of Variance 0.3582 0.1888 0.146 0.09147 0.08193 0.0598 0.04601
## Cumulative Proportion 0.3582 0.5469 0.693 0.78445 0.86638 0.9262 0.97219
##
## Standard deviation      PC8      PC9
## Standard deviation      0.3612 0.34612
## Proportion of Variance 0.0145 0.01331
## Cumulative Proportion 0.9867 1.00000

pcaCharts <- function(x) {
  x.var <- x$sdev ^ 2
  x.pvar <- x.var/sum(x.var)
  print("proportions of variance:")
  print(x.pvar)

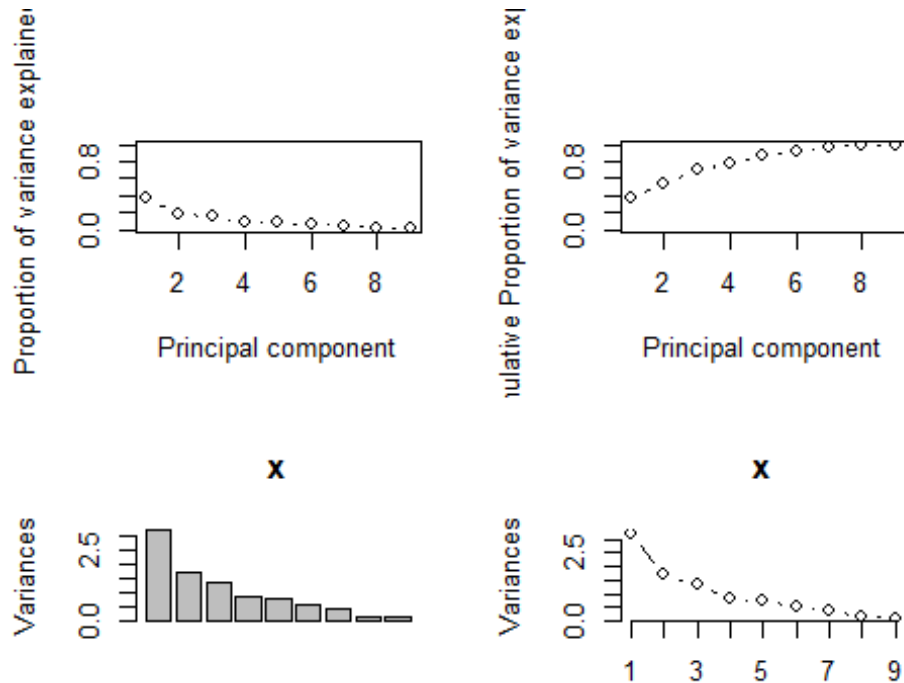
  par(mfrow=c(2,2))
  plot(x.pvar,xlab="Principal component", ylab="Proportion of variance
explained", ylim=c(0,1), type='b')
  plot(cumsum(x.pvar),xlab="Principal component", ylab="Cumulative Proportion
of variance explained", ylim=c(0,1), type='b')
  screeplot(x)
  screeplot(x,type="l")
}
```

```

    par(mfrow=c(1,1))
  }
  pcaCharts(dimPCA)

## [1] "proportions of variance:"
## [1] 0.35815489 0.18877760 0.14604473 0.09147301 0.08192947 0.05979687
0.04601363
## [8] 0.01449860 0.01331121

```



```

###Check VIF factor(Multicollinearity)
library(car)

## Warning: package 'car' was built under R version 4.1.2

## Loading required package: carData

## Warning: package 'carData' was built under R version 4.1.2

#create vector of VIF values
model=lm(y~.,D)
vif_values <- vif(model)

#create horizontal bar chart to display each VIF value
barplot(vif_values, main = "VIF Values", horiz = TRUE, col = "steelblue")
#add vertical line at 5
abline(v = 5, lwd = 3, lty = 2)
##

```

```

library(tidyverse) #tidyverse for easy data manipulation and visualization

## Warning: package 'tidyverse' was built under R version 4.1.1


## -- Attaching packages ----- tidyverse
1.3.1 --

## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.3      v dplyr  1.0.7
## v tidyr   1.1.3      v stringr 1.4.0
## v readr   2.0.0      v forcats 0.5.1

## Warning: package 'ggplot2' was built under R version 4.1.1
## Warning: package 'stringr' was built under R version 4.1.1

## -- Conflicts -----
tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## x dplyr::recode()  masks car::recode()
## x purrr::some()    masks car::some()

```



```

library(caret) #caret for easy machine Learning workflow

## Warning: package 'caret' was built under R version 4.1.2

```

```

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift

library(leaps)      #leaps, for computing stepwise regression

## Warning: package 'leaps' was built under R version 4.1.2

model <- lm(y~.,data=D)
library(MASS)

##
## Attaching package: 'MASS'

## The following object is masked from 'package:dplyr':
##
## select

# Fit the full model
full.model <- lm(y ~., data = D)
# Stepwise regression model
step.model <- stepAIC(full.model, direction = "both",
                      trace = FALSE)
summary(step.model)

##
## Call:
## lm(formula = y ~ x2 + x7 + x8 + x9, data = D)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.3519 -0.5612 -0.0856  0.6972  3.2802
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.8217034  7.7847061  -0.234  0.81705
## x2           0.0038186  0.0007051   5.416 1.67e-05 ***
## x7           0.2168941  0.0886759   2.446  0.02252 *
## x8          -0.0040149  0.0013983  -2.871  0.00863 **
## x9          -0.0016349  0.0012460  -1.312  0.20244
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.681 on 23 degrees of freedom
## Multiple R-squared:  0.8012, Adjusted R-squared:  0.7666
## F-statistic: 23.17 on 4 and 23 DF, p-value: 8.735e-08

```

```
#lm(formula = y ~ x2 + x7 + x8 + x9, data = D) ##Best model using stepwise regression
```

```
smp_siz = floor(0.75*nrow(D))
set.seed(123)
train_ind = sample(seq_len(nrow(D)),size = smp_siz)
train =D[train_ind,]
test=D[-train_ind,]
```

```
##Fitting the important model
```

```
modell1=lm(lm(formula = y ~ x2 + x7 + x8 + x9, data = train))
summary(modell1)
```

```
##
## Call:
## lm(formula = lm(formula = y ~ x2 + x7 + x8 + x9, data = train))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.41776 -0.42020 -0.03068  0.98249  2.57813
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -4.129e+00  7.507e+00  -0.550  0.58992
## x2           3.234e-03  8.151e-04   3.968  0.00110 **
## x7           2.298e-01  9.089e-02   2.529  0.02235 *
## x8          -4.440e-03  1.448e-03  -3.066  0.00739 **
## x9           2.662e-05  1.607e-03   0.017  0.98698
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.496 on 16 degrees of freedom
## Multiple R-squared:  0.8419, Adjusted R-squared:  0.8024
## F-statistic: 21.3 on 4 and 16 DF, p-value: 3.02e-06
```

```
##x2,x7,x8 are three significant variables
```

```
pre=predict(modell1,test)
z=data.frame(test[,1],floor(pre))
z
```

```
##      test...1. floor.pre.
## 1           10          6
## 2           11          8
## 6           11         11
## 12          10          7
## 21           3          6
```

```
## 23      4      4
## 27      2      1
```

```
sigma(model1)/mean(D[,1])
```

```
## [1] 0.2148519
```

```
##21% mean prediction error
```